

# Sentiment Analysis



This report covers the approach taken for sentiment analysis of movie reviews. The report is divided into the following sections:

1. Dataset
2. Data Exploration
3. Classification and Sentiment Labeling

“Proin metus  
urna porta non,  
tincidunt ornare.  
Class aptent  
taciti sociosqu  
ad per inceptos  
hamenaeos.”

-Leo Praesen



## Dataset:

The dataset is comprised of tab-separated files with phrases from the Rotten Tomatoes dataset. The train/test split has been preserved for the purposes of benchmarking, but the sentences have been shuffled from their original order. Each Sentence has been parsed into many phrases by the Stanford parser. Each phrase has a Phraseld. Each sentence has a Sentenceld. Phrases that are repeated (such as short/common words) are only included once in the data.

- train.tsv contains the phrases and their associated sentiment labels. We have additionally provided a Sentenceld so that you can track which phrases belong to a single sentence.
- test.tsv contains just phrases. You must assign a sentiment label to each phrase.

The sentiment labels are:

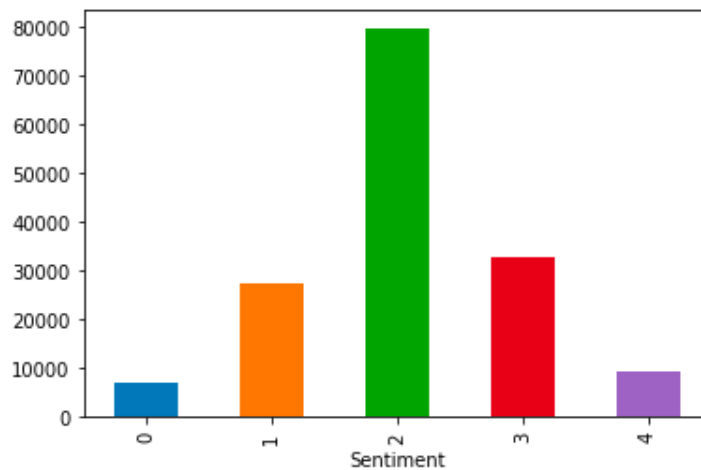
- 0 - negative
- 1 - somewhat negative
- 2 - neutral
- 3 - somewhat positive
- 4 - positive

## Data Exploration:


Data exploration in sentiment analysis involves examining the dataset for the following features:

1. Visualizing the distribution of the dataset:


The dataset distribution is shown in the figure below. It can be observed that the dataset is not balanced and some classes have more values than the others.



2. To balance the dataset we calculate the size of the smallest class and make sure that every class has the same number of samples. As the smallest class has 7,072 samples, we prune the dataset by selecting 7,072 samples from each class. The data distribution after sampling is as follows:



Class 0:



Class 3:



Class 4:



The 5 classes have “movie”, “one” and “movie” as common words that are repeated a lot. We add these words to the list of stop words.

The training set has 12212 unique words and the maximum review size is 48.



The training set has 12212 unique words and the maximum review size is 48.

## Classification and Sentiment Labeling:

For the classification models we start with Machine Learning Approaches and then move on to deep learning methods. The classification approaches are as follows:

### 1. Tf-IDf based classification:

We use Multinomial Naive Bayes and Random Forest Classifiers. The input variable is the Term-Frequency Inverse Document Frequency scores.

The classification reports are as follows:

Multinomial Naive Bayes:

accuracy 0.47219079939668174

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.53      | 0.71   | 0.61     | 2122    |
| 1            | 0.43      | 0.35   | 0.38     | 2185    |
| 2            | 0.46      | 0.22   | 0.29     | 2127    |
| 3            | 0.40      | 0.37   | 0.38     | 2120    |
| 4            | 0.49      | 0.73   | 0.59     | 2054    |
| micro avg    | 0.47      | 0.47   | 0.47     | 10608   |
| macro avg    | 0.46      | 0.47   | 0.45     | 10608   |
| weighted avg | 0.46      | 0.47   | 0.45     | 10608   |

Random Forest Classifier:

accuracy 0.5299773755656109

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.66      | 0.63   | 0.64     | 2122    |
| 1            | 0.46      | 0.35   | 0.40     | 2185    |
| 2            | 0.48      | 0.66   | 0.55     | 2127    |
| 3            | 0.43      | 0.38   | 0.40     | 2120    |
| 4            | 0.63      | 0.64   | 0.63     | 2054    |
| micro avg    | 0.53      | 0.53   | 0.53     | 10608   |
| macro avg    | 0.53      | 0.53   | 0.53     | 10608   |
| weighted avg | 0.53      | 0.53   | 0.52     | 10608   |

## 2. Word2Vec models using the Google News Vectors:

The GoogleNews-vectors through Word2Vec makes use of a pre-trained model that assigns word vectors to the words. These are passed as input to the classifier.

The words are tokenized and then passed through a word vectorizer that uses the pretrained model of GoogleNews vectors.

The classifiers used here are:

Logistic Regression Classifier:

```
accuracy 0.5258295625942685
          precision    recall  f1-score   support

     0         0.55         0.66         0.60         2122
     1         0.45         0.31         0.37         2185
     2         0.54         0.63         0.58         2127
     3         0.44         0.33         0.38         2120
     4         0.58         0.72         0.64         2054

  micro avg         0.53         0.53         0.53        10608
  macro avg         0.51         0.53         0.51        10608
weighted avg         0.51         0.53         0.51        10608
```

Random Forest Classifier:

```
accuracy 0.4671945701357466
          precision    recall  f1-score   support

     0         0.53         0.61         0.57         2122
     1         0.38         0.37         0.37         2185
     2         0.46         0.48         0.47         2127
     3         0.37         0.31         0.34         2120
     4         0.59         0.56         0.58         2054

  micro avg         0.47         0.47         0.47        10608
  macro avg         0.46         0.47         0.47        10608
weighted avg         0.46         0.47         0.46        10608
```



### 3. Doc2Vec to classify the data at a document level using DBOW (Distributed Bag of Words)

Logistic Regression Classifier:

```
accuracy 0.4549396681749623
      precision    recall  f1-score   support

     0         0.52      0.60      0.56      2103
     1         0.35      0.23      0.27      2143
     2         0.41      0.57      0.48      2102
     3         0.37      0.25      0.30      2112
     4         0.55      0.63      0.59      2148

 micro avg      0.45      0.45      0.45     10608
 macro avg      0.44      0.46      0.44     10608
weighted avg      0.44      0.45      0.44     10608
```

Random Forest Classifier:

```
accuracy 0.4549396681749623
      precision    recall  f1-score   support

     0         0.52      0.60      0.56      2103
     1         0.35      0.23      0.27      2143
     2         0.41      0.57      0.48      2102
     3         0.37      0.25      0.30      2112
     4         0.55      0.63      0.59      2148

 micro avg      0.45      0.45      0.45     10608
 macro avg      0.44      0.46      0.44     10608
weighted avg      0.44      0.45      0.44     10608
```

It is observed that classifying at a document level is not too effective.

### 4. Deep Learning models:

The advantage of using Deep Learning is that we can use models that can capture more word combinations than n-gram models and the use of RNNs and LSTMs allow capturing sequential information. The models primarily follow the following preprocessing steps:

1. Tokenization
2. Embedding the tokens into an embedding matrix
3. Train the model
4. Make predictions

## LSTM Model

The LSTM model is described as follows:

Model: "sequential\_24"

| Layer (type)                | Output Shape      | Param # |
|-----------------------------|-------------------|---------|
| =====                       |                   |         |
| embedding_7 (Embedding)     | (None, None, 100) | 1374000 |
| =====                       |                   |         |
| lstm_13 (LSTM)              | (None, None, 64)  | 42240   |
| =====                       |                   |         |
| lstm_14 (LSTM)              | (None, 32)        | 12416   |
| =====                       |                   |         |
| dense_68 (Dense)            | (None, 5)         | 165     |
| =====                       |                   |         |
| Total params: 1,428,821     |                   |         |
| Trainable params: 1,428,821 |                   |         |
| Non-trainable params: 0     |                   |         |
| =====                       |                   |         |

The results have the following accuracy: 66.02

## CNN model

The CNN models is as follows:

Model: "sequential\_29"

| Layer (type)                 | Output Shape    | Param # |
|------------------------------|-----------------|---------|
| =====                        |                 |         |
| embedding_12 (Embedding)     | (None, 48, 100) | 1374000 |
| =====                        |                 |         |
| dropout_33 (Dropout)         | (None, 48, 100) | 0       |
| =====                        |                 |         |
| conv1d_5 (Conv1D)            | (None, 48, 64)  | 19264   |
| =====                        |                 |         |
| global_max_pooling1d_5 (Glob | (None, 64)      | 0       |
| =====                        |                 |         |
| dense_77 (Dense)             | (None, 128)     | 8320    |
| =====                        |                 |         |
| dropout_34 (Dropout)         | (None, 128)     | 0       |
| =====                        |                 |         |
| dense_78 (Dense)             | (None, 5)       | 645     |
| =====                        |                 |         |
| Total params: 1,402,229      |                 |         |
| Trainable params: 1,402,229  |                 |         |
| Non-trainable params: 0      |                 |         |
| =====                        |                 |         |

The accuracy is as follows: 66.84

## CNN+CRU model

The CNN+GRU model is as follows:

Model: "sequential\_30"

| Layer (type)                   | Output Shape    | Param # |
|--------------------------------|-----------------|---------|
| =====                          |                 |         |
| embedding_13 (Embedding)       | (None, 48, 100) | 1374000 |
| -----                          |                 |         |
| conv1d_6 (Conv1D)              | (None, 48, 64)  | 19264   |
| -----                          |                 |         |
| max_pooling1d_1 (MaxPooling1D) | (None, 24, 64)  | 0       |
| -----                          |                 |         |
| dropout_35 (Dropout)           | (None, 24, 64)  | 0       |
| -----                          |                 |         |
| gru_1 (GRU)                    | (None, 24, 128) | 74112   |
| -----                          |                 |         |
| dropout_36 (Dropout)           | (None, 24, 128) | 0       |
| -----                          |                 |         |
| flatten_1 (Flatten)            | (None, 3072)    | 0       |
| -----                          |                 |         |
| dense_79 (Dense)               | (None, 128)     | 393344  |
| -----                          |                 |         |
| dropout_37 (Dropout)           | (None, 128)     | 0       |
| -----                          |                 |         |
| dense_80 (Dense)               | (None, 5)       | 645     |
| =====                          |                 |         |
| Total params: 1,861,365        |                 |         |
| Trainable params: 1,861,365    |                 |         |
| Non-trainable params: 0        |                 |         |
| -----                          |                 |         |

The accuracy is as follows: 67.08

## Bi-Directional GRU model

The bidirectional GRU model is as follows:

Model: "sequential\_31"

| Layer (type)                 | Output Shape    | Param # |
|------------------------------|-----------------|---------|
| =====                        |                 |         |
| embedding_14 (Embedding)     | (None, 48, 100) | 1374000 |
| -----                        |                 |         |
| spatial_dropout1d_1 (Spatial | (None, 48, 100) | 0       |
| -----                        |                 |         |
| bidirectional_1 (Bidirection | (None, 256)     | 175872  |
| -----                        |                 |         |
| dropout_38 (Dropout)         | (None, 256)     | 0       |
| -----                        |                 |         |
| dense_81 (Dense)             | (None, 5)       | 1285    |
| =====                        |                 |         |
| Total params: 1,551,157      |                 |         |
| Trainable params: 1,551,157  |                 |         |
| Non-trainable params: 0      |                 |         |
| -----                        |                 |         |

The accuracy is as follows: 66.86

## Using GloVe word embeddings:

In this model we use pre-trained word embeddings. The pre-trained model here is called GloVe - Global Vectors for Word Representation and it allows us to have the embedding matrix with pre-trained weights for the words.

The model used is as follows:

Model: "sequential\_34"

| Layer (type)                 | Output Shape    | Param # |
|------------------------------|-----------------|---------|
| embedding_16 (Embedding)     | (None, 48, 100) | 1374000 |
| spatial_dropout1d_3 (Spatial | (None, 48, 100) | 0       |
| bidirectional_4 (Bidirection | (None, 48, 256) | 175872  |
| bidirectional_5 (Bidirection | (None, 128)     | 123264  |
| dropout_40 (Dropout)         | (None, 128)     | 0       |
| dense_83 (Dense)             | (None, 5)       | 645     |
| Total params: 1,673,781      |                 |         |
| Trainable params: 1,673,781  |                 |         |
| Non-trainable params: 0      |                 |         |

The accuracy is as follows: 68.49